

Xuhao Luo

201 N Goodwin Ave, Urbana, IL, 61801
xuhaol2@illinois.edu ◊ LinkedIn

Education

University of Illinois Urbana-Champaign

Ph.D. Candidate in Computer Science

University of California San Diego

M.S. in Computer Science

University of Science and Technology of China

B.S. in Applied Physics

Aug. 2021 - Nov. 2026 (Anticipated)

Sep. 2019 - Mar. 2021

Sep. 2015 - Jun. 2019

Interests

Distributed Systems, Memory Systems, and Storage Systems for AI

Research Publication

- **Xuhao Luo**, Emaan Atique, Aishwarya Ganesan, Ramnatthan Alagappan, **SkyLog: Reducing Cost and Latency in a Multi-Cloud Shared Log** (In submission to NSDI)
- **Xuhao Luo**, Shreesha G. Bhat*, Jiyu Hu*(equal contribution), Aishwarya Ganesan, Ramnatthan Alagappan, **LazyLog: A New Shared Log Abstraction for Low-Latency Applications (SOSP 2024)** *Best Paper Award, Invited to ACM ToCS*
- **Xuhao Luo**, Ramnatthan Alagappan, Aishwarya Ganesan, **SplitFT: Fault Tolerance for Disaggregated Datacenters via Remote Memory Logging (EuroSys 2024)**
- **Xuhao Luo**, Weihai Shen, Shuai Mu, Tianyin Xu, **DepFast: Orchestrating Code of Quorum Systems (USENIX ATC 2022)**
- Shreesha G. Bhat, Tony Hong, **Xuhao Luo**, Jiyu Hu, Aishwarya Ganesan, Ramnatthan Alagappan, **Low End-to-End Latency atop a Speculative Shared Log with Fix-Ante Ordering (OSDI 2025)**
- Zhiyuan Guo*, Yizhou Shan*(co-first author), **Xuhao Luo**, Yutong Huang, Yiyang Zhang, **Clio: A Hardware-Software Co-Designed Disaggregated Memory System (ASPLOS 2022)**

Experience

Google

Software Engineer Intern, Host: Xiaozhou (Steve) Li

May. 2026 - Aug. 2026

Sunnyvale, CA, USA

- Designed and implemented a general framework to extend Bigtable compaction with user-defined functions. Deployed in production and demonstrated its feasibility to replace certain data processing pipeline.

Amazon Web Service

Applied Scientist Intern, Mentor: Prof. George Amvrosiadis, Visiting Scholar

May. 2024 - Aug. 2024

Seattle, WA, USA

- Investigated and improved ShardStore (storage system used by AWS S3) reclamation, and built a tool to evaluate different reclamation policy.

Amazon Web Service

Applied Scientist Intern, Mentor: Shen Li, Principle Engineer

May. 2022 - Aug. 2022

Seattle, WA, USA

- Improved the performance and reliability of the volume metadata updating workflow for AWS S3 volume metadata cache service with Amazon Quantum Ledger Database (QLDB).

University of Illinois Urbana-Champaign

Research Assistant

May. 2021 - Now

Urbana, IL, USA

- *SkyLog*: Built a shared log abstraction for multicloud applications, with seamless reconfiguration and reduced egress cost.
- *LazyLog*: Built a new shared log system that offers low-latency appends for latency-critical applications by lazily ordering log entries across multiple storage shards.
- *SplitFT*: Built a new fault-tolerant approach for storage-centric cloud applications by replicating WALs on remote nodes using RDMA.
- *DepFast*: Built a framework using C++ coroutine to implement and reason about fail-slow tolerant distributed systems in an easy and effective way.

Microsoft Research

Research Intern

Jun. 2020 - Sep. 2020

Beijing, China

- Designed and implemented a task scheduling and dispatching system for distributed ML, and a CUDA-based high-performance inter-GPU communication channel for distributed ML.
- Multi-GPU collective operation(AllReduce, AllGather, Broadcast) throughput outperforms NCCL by at most 18.4%.

University of California San Diego
Research Assistant, advised by Prof. Yiyang Zhang

Sep. 2019 - Dec. 2020
La Jolla, CA, USA

- Implemented a hardware/software co-designed disaggregated memory with FPGA. Worked on the network stack design and implementation.

University of Science and Technology of China
Graduation Project, Advisor: Prof. Xi Jin

Jan. 2019 - May. 2019

- Designed and implemented a general 8-bit quantized convolution module on Xilinx Virtex FPGA.
- Developed the TensorFlow C++API for the hardware accelerator using OpenCL. On ResNet-50 CNN, achieved a speedup of 5.17x and a memory usage reduction of 66% compared with the CPU TensorFlow implementation on Xeon E5 2686.

Skills

Language	C/C++, Python, Go, Java, Rust, Haskell, Verilog
Tools/Framework	RDMA, vLLM, SGLang, TensorFlow, Docker, Zookeeper, LLVM

Honors and Awards

- | | |
|--|----------|
| • SOSP'24 Best Paper Award | Nov 2024 |
| • SOSP'24 Student Travel Grant | Nov 2024 |
| • EuroSys'24 Student Travel Grant | Apr 2024 |
| • ASPLOS'22 Student Travel Grant | Feb 2022 |
| • USTC Class of 2019 Outstanding Graduates | May 2019 |

Services

- EuroSys'26 Shadow PC
- OSDI'23 Artifact Evaluation Committee
- USENIX ATC'23 Artifact Evaluation Committee
- ACM TAAS (Transactions on Autonomous and Adaptive Systems) Invited Reviewer